

R-package **FME** : inverse modelling, sensitivity, Monte Carlo - applied to a steady-state model

Karline Soetaert
NIOO-CEME
The Netherlands

Abstract

Rpackage **FME** (?) contains functions for model calibration, sensitivity, identifiability, and Monte Carlo analysis of nonlinear models.

This vignette, (`vignette("FMEsteady")`), applies **FME** to a partial differential equation, solved with a steady-state solver from package **rootSolve**

A similar vignette (`vignette("FMEdyna")`), applies the functions to a dynamic simulation model, solved with integration routines from package **deSolve**

A third vignette (`vignette("FMEother")`), applies the functions to a simple nonlinear model

`vignette("FMEcmc")` tests the Markov chain Monte Carlo (MCMC) implementation

Keywords: steady-state models, differential equations, fitting, sensitivity, Monte Carlo, identifiability, R.

1. A steady-state model of oxygen in a marine sediment

This is a simple model of oxygen in a marine (submersed) sediment, diffusing along a spatial gradient, with imposed upper boundary concentration oxygen is consumed at maximal fixed rate, and including a monod limitation.

See (?) for a description of reaction-transport models.

The constitutive equations are:

$$\begin{aligned}\frac{\partial O_2}{\partial t} &= -\frac{\partial Flux}{\partial x} - cons \cdot \frac{O_2}{O_2 + k_s} \\ Flux &= -D \cdot \frac{\partial O_2}{\partial x} \\ O_2(x=0) &= upO2\end{aligned}$$

```
> par(mfrow=c(2,2))  
> require(FME)
```

First the model parameters are defined...

```
> pars <- c(upO2=360, # concentration at upper boundary, mmolO2/m3  
+          cons=80,   # consumption rate, mmolO2/m3/day
```

R-package **FME** : inverse modelling, sensitivity, Monte Carlo - applied to a steady-state model

```
+          ks=1,          # O2 half-saturation ct, mmolO2/m3
+          D=1)          # diffusion coefficient, cm2/d
```

Next the sediment is vertically subdivided into 100 grid cells, each 0.05 cm thick.

```
> n <- 100                      # nr grid points
> dx <- 0.05    #cm
> dX <- c(dx/2,rep(dx,n-1),dx/2) # dispersion distances; half dx near boundaries
> X  <- seq(dx/2,len=n,by=dx)    # distance from upper interface at middle of box
```

The model function takes as input the parameter values and returns the steady-state condition of oxygen. Function `steady.band` from package **rootSolve** ((?)) does this in a very efficient way (see (?)).

```
> O2fun <- function(pars)
+ {
+   derivs<-function(t,O2,pars)
+   {
+     with (as.list(pars),{
+       Flux <- -D* diff(c(upO2,O2,O2[n]))/dX
+       dO2  <- -diff(Flux)/dx-cons*O2/(O2+ks)
+       return(list(dO2,UpFlux = Flux[1],LowFlux = Flux[n+1]))
+     })
+   }
+   # Solve the steady-state conditions of the model
+   ox <- steady.band(y=runif(n),func=derivs,parms=pars,nspec=1,positive=TRUE)
+   data.frame(X=X,O2=ox$y)
+ }
```

The model is run

```
> ox<-O2fun(pars)
```

and the results plotted...

```
> plot(ox$O2,ox$X,ylim=rev(range(X)),xlab="mmol/m3",
+       main="Oxygen", ylab="depth, cm",type="l",lwd=2)
```

2. Global sensitivity analysis : Sensitivity ranges

The sensitivity of the oxygen profile to parameter `cons`, the consumption rate is estimated. We assume a normally distributed parameter, with mean = 80 (`parMean`), and a variance=100 (`parCovar`). The model is run 100 times (`num`).

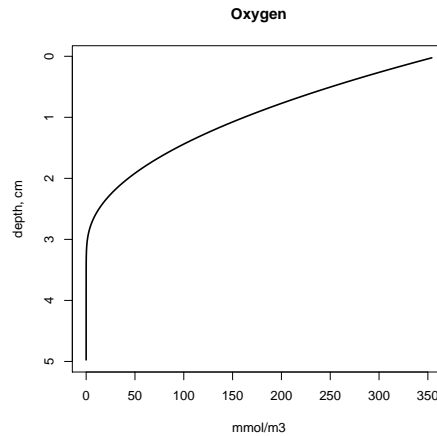


Figure 1: The modeled oxygen profile - see text for R-code

```
> print(system.time(
+ Sens2 <- sensRange(parms=pars,func=O2fun,dist="norm",
+                     num=100,parMean=c(cons=80),parCovar=100)
+ ))
```

```
user  system elapsed
0.97   0.00   0.97
```

The results can be plotted in two ways:

```
> par(mfrow=c(1,2))
> plot(Sens2,xyswap=TRUE,xlab= "O2",
+       ylab="depth, cm",main="Sensitivity runs")
> plot(summary(Sens2),xyswap=TRUE,xlab= "O2",
+       ylab="depth, cm",main="Sensitivity ranges")
> par(mfrow=c(1,1))
```

3. Local sensitivity analysis : Sensitivity functions

Local sensitivity analysis starts by calculating the sensitivity functions

```
> O2sens <- sensFun(func=O2fun,parms=pars)
```

The summary of these functions gives information about which parameters have the largest effect (univariate sensitivity):

```
> summary(O2sens)
```




Figure 2: Results of the sensitivity run - left: all model runs, right: summary - see text for R-code

	value	scale	L1	L2	Mean	Min	Max	N
upO2	360	360	7.0	0.88	7.0	1.0e+00	13.4176	100
cons	80	80	8.5	1.20	-8.5	-2.3e+01	-0.0084	100
ks	1	1	2.2	0.37	2.2	1.2e-04	9.6137	100
D	1	1	8.1	1.14	8.1	8.4e-03	22.0312	100

In bivariate sensitivity the pair-wise relationship and the correlation is estimated and/or plotted:

```
> pairs(O2sens)
```

```
> cor(O2sens[,-(1:2)])
```

	upO2	cons	ks	D
upO2	1.0000000	-0.9781166	0.8375806	0.9787945
cons	-0.9781166	1.0000000	-0.9326397	-0.9998910
ks	0.8375806	-0.9326397	1.0000000	0.9317287
D	0.9787945	-0.9998910	0.9317287	1.0000000

Multivariate sensitivity is done by estimating the collinearity between parameter sets (?).

```
> Coll <- collin(O2sens)
```

```
> Coll
```

	upO2	cons	ks	D	N	collinearity
1	0	0	1	1	2	4.4
2	0	1	1	1	3	128.6
3	1	1	1	1	4	143.7




Figure 3: pairs plot - see text for R-code




Figure 4: collinearity - see text for R-code

R-package **FME** : inverse modelling, sensitivity, Monte Carlo - applied to a steady-state model

```
> plot(Coll,log="y")
```

4. Fitting the model to the data

Assume both the oxygen flux at the upper interface and a vertical profile of oxygen has been measured.

These are the data:

```
> O2dat <- data.frame(x=seq(0.1,3.5,by=0.1),
+   y = c(279,260,256,220,200,203,189,179,165,140,138,127,116,
+   109,92,87,78,72,62,55,49,43,35,32,27,20,15,15,10,8,5,3,2,1,0))
> O2depth <- cbind(name="O2",O2dat)          # oxygen versus depth
> O2flux <- c(UpFlux=170)                    # measured flux
```

First a function is defined that returns only the required model output.

```
> O2fun2 <- function(pars)
+ {
+   derivs<-function(t,O2,pars)
+   {
+     with (as.list(pars),{
+       Flux <- -D*diff(c(upO2,O2,O2[n]))/dX
+       dO2 <- -diff(Flux)/dx-cons*O2/(O2+ks)
+       return(list(dO2,UpFlux = Flux[1],LowFlux = Flux[n+1]))
+     })
+   }
+   ox <- steady.band(y=runif(n),func=derivs,parms=pars,nspec=1,
+     positive=TRUE,rtol=1e-8,atol=1e-10)
+   list(data.frame(x=X,O2=ox$y),
+     UpFlux=ox$UpFlux)
+ }
```

The function used in the fitting algorithm returns an instance of type `modCost`. This is created by calling function `modCost` twice. First with the modeled oxygen profile, then with the modeled flux.

```
> Objective <- function (P)
+ {
+   Pars <- pars
+   Pars[names(P)]<-P
+   modO2 <- O2fun2(Pars)
+ }
```

```

+ # Model cost: first the oxygen profile
+ Cost <- modCost(obs=O2depth,model=modO2[[1]],x="x",y="y")
+
+ # then the flux
+ modFl <- c(UpFlux=modO2$UpFlux)
+ Cost <- modCost(obs=O2flux,model=modFl,x=NULL,cost=Cost)
+
+ return(Cost)
+ }

```

We first estimate the identifiability of the parameters, given the data:

```

> print(system.time(
+ sF<-sensFun(Objective,parms=pars)
+ ))

```

```

      user  system elapsed
0.11      0.00      0.11

```

```

> summary(sF)

```

	value	scale	L1	L2	Mean	Min	Max	N
up02	360	360	4.25	0.97	4.25	0.5069	13.3	36
cons	80	80	3.68	0.99	-3.65	-15.3722	0.5	36
ks	1	1	0.40	0.14	0.40	-0.0069	3.1	36
D	1	1	3.68	0.99	3.68	0.0342	15.4	36

```

> collin(sF)

```

	up02	cons	ks	D	N	collinearity
1	0	0	1	1	2	4.2
2	0	1	1	1	3	50.6
3	1	1	1	1	4	51.0

The collinearity of the full set is too high, but as the oxygen diffusion coefficient is well known, it is left out of the fitting. The combination of the three remaining parameters has a low enough collinearity to enable automatic fitting. The parameters are constrained to be >0

```

> collin(sF,parset=c("up02","cons","ks"))

```

	up02	cons	ks	D	N	collinearity
1	1	1	1	0	3	14

```

> print(system.time(
+ Fit<-modFit(p=c(up02=360,cons=80,ks=1),
+               f=Objective,lower=c(0,0,0))
+ ))

```




Figure 5: residuals - see text for R-code

```

user  system elapsed
0.92   0.00   0.93

> (SFit<-summary(Fit))

Parameters:
      Estimate Std. Error t value Pr(>|t|)
up02  292.937     2.104 139.241  <2e-16 ***
cons   49.686     2.367  20.991  <2e-16 ***
ks      1.297     1.363   0.951   0.348
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.401 on 33 degrees of freedom

Parameter correlation:
      up02  cons  ks
up02 1.0000 0.5791 0.2976
cons 0.5791 1.0000 0.9013
ks   0.2976 0.9013 1.0000

We next plot the residuals

> plot(Objective(Fit$par),xlab="depth",ylab="",main="residual",legpos="top")

and show the best-fit model

> Pars <- pars
> Pars[names(Fit$par)]<- Fit$par
> mod02 <- 02fun(Pars)

```





Figure 6: Best fit model - see text for R-code

```
> plot(O2depth$y,O2depth$x,ylim=rev(range(O2depth$x)),pch=18,
+      main="Oxygen-fitted", xlab="mmol/m3",ylab="depth, cm")
> lines(modO2$O2,modO2$X)
```

5. Running a Markov chain Monte Carlo

We use the parameter covariances of previous fit to update parameters, while the mean squared residual of the fit is use as prior fo the model variance.

```
> Covar    <- SFit$cov.scaled * 2.4^2/3
> s2prior  <- SFit$modVariance
```

We run an adaptive Metropolis, making sure that ks does not become negative...

```
> print(system.time(
+ MCMC <- modMCMC(f=Objective,p=Fit$par,jump=Covar,niter=1000,ntrydr=2,
+               var0=s2prior,wvar0=1,updatecov=100,lower=c(NA,NA,0))
+ ))
```

number of accepted runs: 718 out of 1000 (71.8%)

```
  user  system elapsed
31.00   0.00   31.34
```

```
> MCMC$count
```

```
dr_steps    Alfasteps  num_accepted num_covupdate
      683         2049         718           9
```

Plotting the results is similar to previous cases.

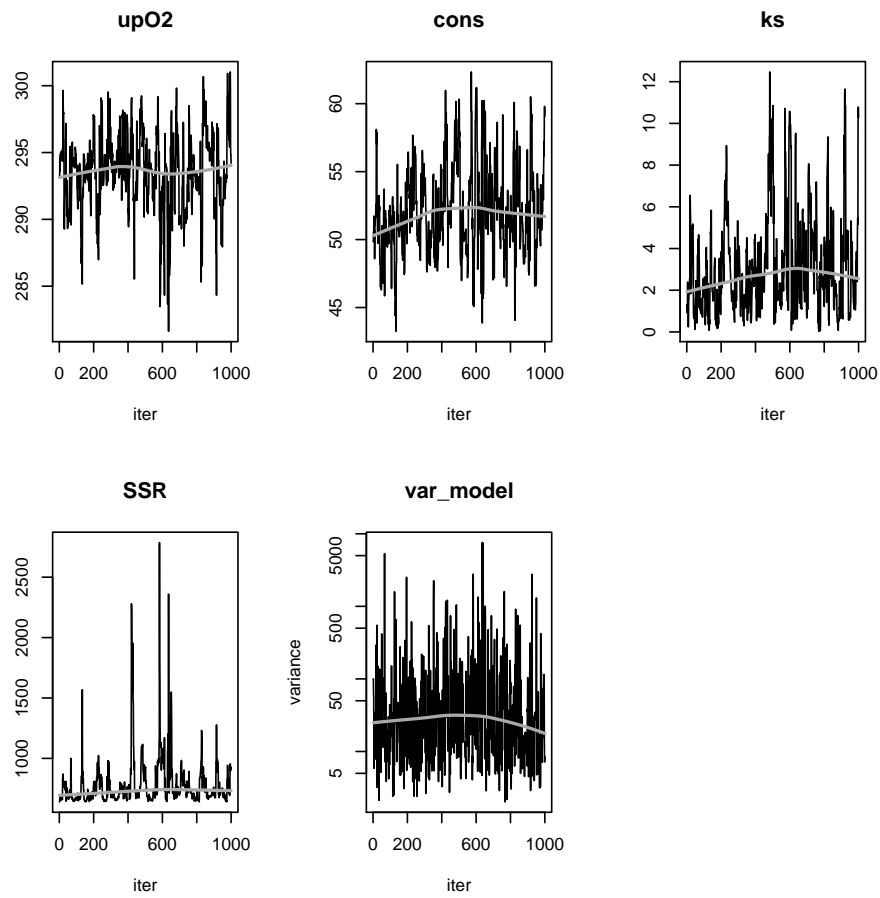


Figure 7: MCMC plot results - see text for R-code

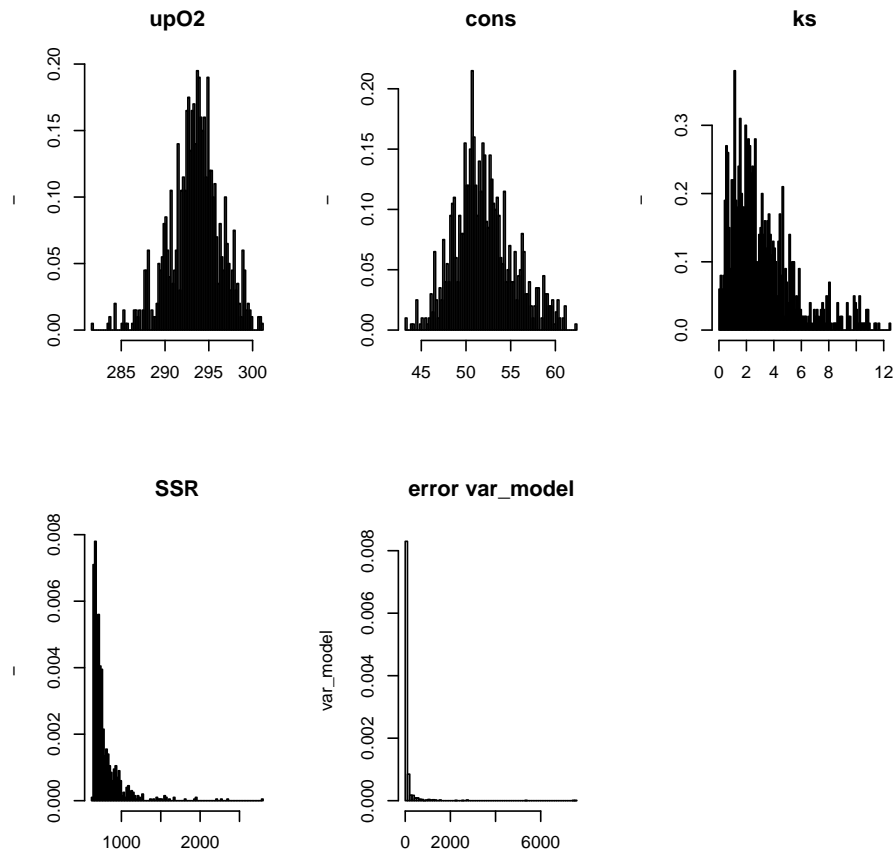


Figure 8: MCMC histogram results - see text for R-code

```
> plot(MCMC, Full=TRUE)
```

```
> hist(MCMC, Full=TRUE)
```

```
> pairs(MCMC, Full=TRUE)
```

or summaries can be created:

```
> summary(MCMC)
```

	upO2	cons	ks	var_model
mean	293.484639	52.051700	3.14473519	108.301175
sd	2.925883	3.315406	2.37077525	439.738658
min	281.623688	43.251055	0.03086142	2.018593
max	301.030933	62.336057	12.46793442	7567.101202

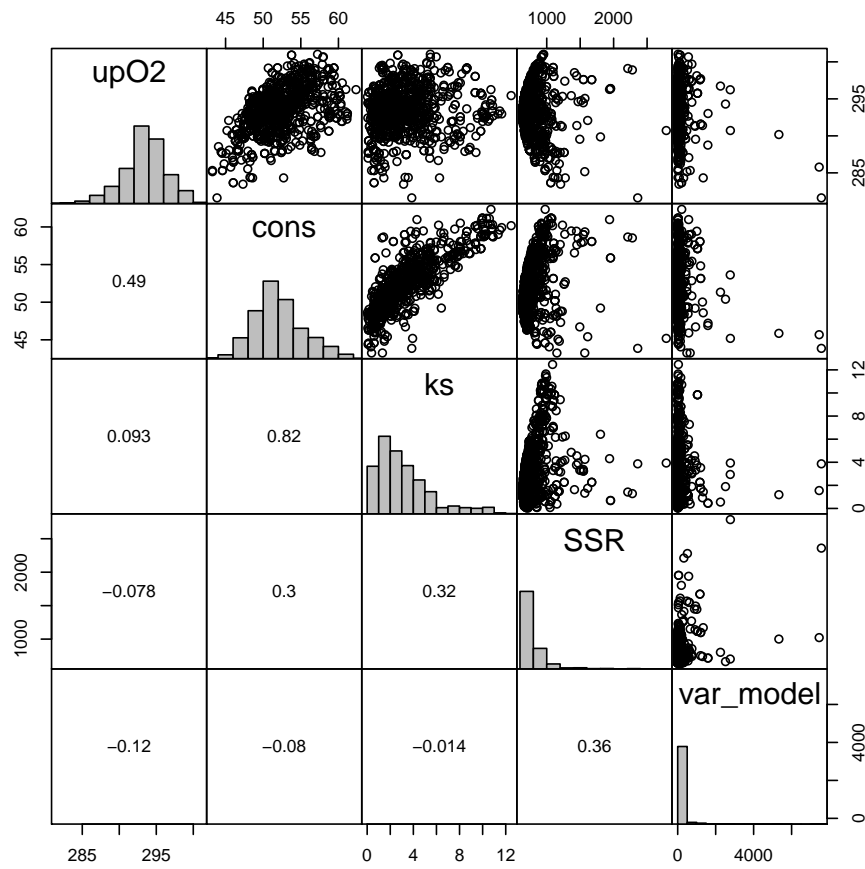


Figure 9: MCMC pairs plot - see text for R-code

```
q025 291.917109 49.939665 1.44063531 10.748082
q050 293.672678 51.736995 2.46861110 24.625373
q075 295.209759 53.878405 4.25521980 62.505279
```

```
> cor(MCMC$pars)
```

```
      up02      cons      ks
up02 1.00000000 0.4891091 0.09315723
cons 0.48910906 1.0000000 0.81871857
ks    0.09315723 0.8187186 1.00000000
```

Note: we pass to `sensRange` the full parameter vector (`parms`) and the parameters sampled during the MCMC (`parInput`).

```
> plot(summary(sensRange(parms=pars,parInput=MCMC$par,f=02fun,num=500)),
+       xyswap=TRUE)
> points(02depth$y,02depth$x)
```

6. Finally

This vignette is made with Sweave (?).

Affiliation:

Karline Soetaert
 Centre for Estuarine and Marine Ecology (CEME)
 Netherlands Institute of Ecology (NIOO)
 4401 NT Yerseke, Netherlands
 E-mail: k.soetaert@nioo.knaw.nl
 URL: <http://www.nioo.knaw.nl/users/ksoetaert>




Figure 10: MCMC range plot - see text for R-code