

R-package **FME** : tests

Karline Soetaert
NIOO-CEME
The Netherlands

Abstract

This vignette tests several applications of Rpackage **FME** .

Keywords: simulation models, differential equations, fitting, sensitivity, monte carlo, identifiability R.

1. Introduction

Here some functions from R-package **FME** are tested.

- Three tests of the implemented MCMC method:
 - The "banana" function, sampling from a curvilinear function (([Laine 2008](#)))
 - A monod function, fitted to a data-series (([Laine 2008](#)))
 - A simple chemical model, fitted to a data series (([Haario, Laine, Mira, and Saksman 2006](#)))
- sensitivity analysis, fitting, MCMC, ... of a simple model of sedimentary oxygen

2. The banana

2.1. the model

This example is from [Laine \(2008\)](#).

A banana-shaped function is created by distorting a two-dimensional Gaussian distribution, with mean = 0 and a covariance matrix τ with unity variances and covariance (of 0.9).

$$\tau = \begin{bmatrix} 0.9 & 0 \\ 0 & 0.9 \end{bmatrix}$$

The distortion is along the second-axis only and given by:

$$\begin{aligned} y_1 &= x_1 \\ y_2 &= x_2 + x_1^2 + 1 \end{aligned}$$

2.2. R-implementation

First the banana function is defined.

```
> Banana <- function (x1,x2)
+ {
+   return(x2 - (x1^2+1))
+ }
```

We also need a function that estimates the probability of a multnormally distributed vector

```
> pmultinorm <- function(vec,mean,Cov)
+
+ {
+   diff <- vec - mean
+   ex   <- -0.5*t(diff) %*% solve(Cov) %*% diff
+   rdet  <- sqrt(det(Cov))
+   power <- -length(diff)*0.5
+   return((2.*pi)^power / rdet * exp(ex))
+ }
```

The target function returns $-2 \cdot \log$ (probability) of the value

```
> BananaSS <- function (p)
+ {
+   P <- c(p[1],Banana(p[1],p[2]))
+   Cov <- matrix(nr=2,data=c(1,0.9,0.9,1))
+   -2*sum(log(pmultinorm(P,mean=0,Cov=Cov)))
+ }
```

The initial proposal covariance (`jump`) is the identity matrix with a variance of 5. The simulated chain is of length 1000 (`niter`). The `modMCMC` function prints the % of accepted runs. More information is in item `count` of its return element.

The First Markov chain is generated with the simple metropolis hastings (MH) algorithm

```
> MCMC <- modMCMC(f=BananaSS, p=c(0,0.5), jump=diag(nrow=2,x=5),
+               niter=1000)
```

number of accepted runs: 109 out of 1000 (10.9%)

```
> MCMC$count
```

dr_steps	Alfasteps	num_accepted	num_covupdate
0	0	109	1

Next we use the adaptive metropolis (AM) algorithm and update the proposal every 100 runs (`updatecov`)

```
> MCMC2 <- modMCMC(f=BananaSS, p=c(0,0.5), jump=diag(nrow=2,x=5),
+                 updatecov=100,niter=1000)
```

number of accepted runs: 175 out of 1000 (17.5%)

```
> MCMC2$count
```

dr_steps	Alfasteps	num_accepted	num_covupdate
0	0	175	10

Then the metropolis algorithm with delayed rejection (DR) is applied; upon rejection one next parameter candidate is tried (`ntrydr`). (note `ntrydr=1` means no delayed rejection steps).

```
> MCMC3 <- modMCMC(f=BananaSS, p=c(0,0.5), jump=diag(nrow=2,x=5),
+                 ntrydr=2,niter=1000)
```

number of accepted runs: 536 out of 1000 (53.6%)

```
> MCMC3$count
```

dr_steps	Alfasteps	num_accepted	num_covupdate
904	2712	536	1

Finally the adaptive metropolis with delayed rejection (DRAM) is used. (Here we also estimate the elapsed CPU time - `print(system.time())` does this)

```
> print(system.time(
+ MCMC4 <- modMCMC(f=BananaSS, p=c(0,0.5), jump=diag(nrow=2,x=5),
+                 updatecov=100,ntrydr=2,niter=1000)
+ ))
```

number of accepted runs: 741 out of 1000 (74.1%)

user	system	elapsed
0.86	0.00	0.85

```
> MCMC4$count
```

dr_steps	Alfasteps	num_accepted	num_covupdate
776	2328	741	10

We plot the generated chains for both parameters and for the four runs in one plot. Calling `plot` with `mfrow=NULL` prevents the plotting function to overrule these settings.

```
> par(mfrow=c(4,2))
> par(mar=c(2,2,4,2))
> plot(MCMC ,mfrow=NULL,main="MH")
> plot(MCMC2,mfrow=NULL,main="AM")
> plot(MCMC3,mfrow=NULL,main="DR")
> plot(MCMC4,mfrow=NULL,main="DRAM")
> mtext(outer=TRUE,side=3,line=-2,at=c(0.05,0.95),c("y1","y2"),cex=1.25)
> par(mar=c(5.1,4.1,4.1,2.1))
```

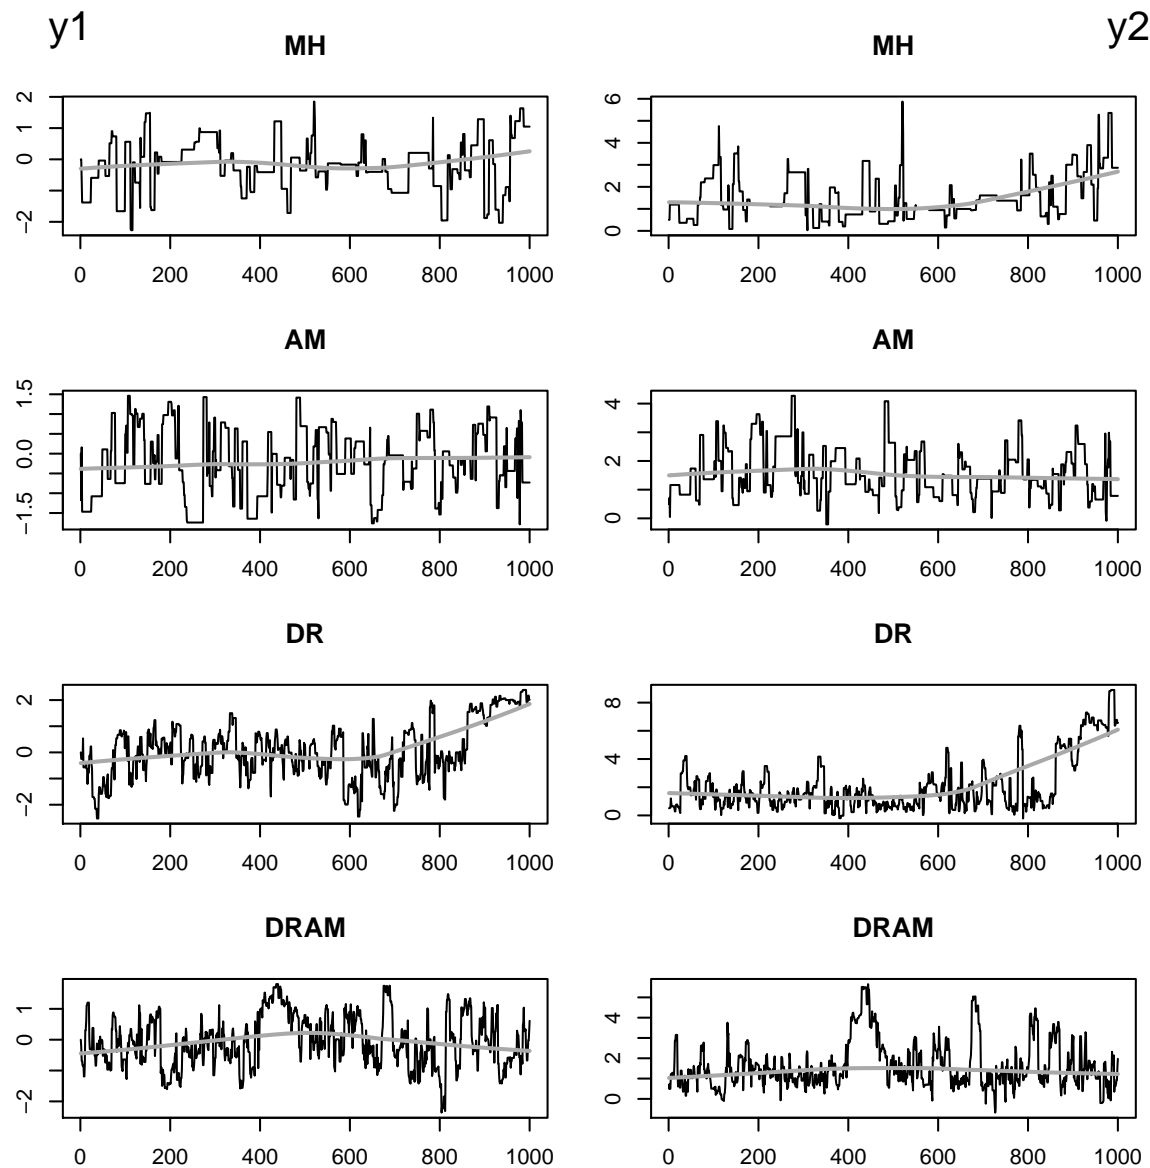


Figure 1: The MCMC chains for the four methods - see text for R-code




Figure 2: The bananas - see text for R-code

The 2-D plots show the banana shape:

```
> par(mfrow=c(2,2))
> xl <- c(-3,3)
> yl <- c(-1,8)
> plot(MCMC$pars,main="MH",xlim=xl,ylim=yl)
> plot(MCMC2$pars,main="AM",xlim=xl,ylim=yl)
> plot(MCMC3$pars,main="DR",xlim=xl,ylim=yl)
> plot(MCMC4$pars,main="DRAM",xlim=xl,ylim=yl)
```

Finally, we test convergence to the original distribution. This can best be done by estimating means and covariances of the transformed parameter values.

```
> trans <- cbind(MCMC4$pars[,1],Banana(MCMC4$pars[,1],MCMC4$pars[,2]))
> colMeans(trans) # was:c(0,0)
```

```
[1] -0.05142509 -0.04655232
```

```
> sd(trans) # was:1
```

```
[1] 0.8242180 0.8740344
```

```
> cor(trans) # 0.9 off-diagonal
```

```

      [,1]      [,2]
[1,] 1.0000000 0.8730663
[2,] 0.8730663 1.0000000

```

3. Fitting a Monod function

3.1. the model

The second example is also discussed in ([Laine 2008](#)) (who quotes Berthouex and Brown, 2002. Statistics for environmental engineers, CRC Press).

The following model:

$$y = \theta_1 \cdot \frac{x}{x + \theta_2} + \epsilon$$

$$\epsilon \sim N(0, I\sigma^2)$$

is fitted to data.

3.2. implementation in R

```
> require(FME)
```

First we input the observations

```
> Obs <- data.frame(x=c( 28, 55, 83, 110, 138, 225, 375), # mg COD/l
+                   y=c(0.053,0.06,0.112,0.105,0.099,0.122,0.125)) # 1/hour
```

The Monod model returns a data.frame, with elements x and y :

```
> Model <- function(p,x) return(data.frame(x=x,y=p[1]*x/(x+p[2])))
```

We first fit the model to the data.

Function `Residuals` estimates the deviances of model versus the data.

```
> Residuals <- function(p) (Obs$y-Model(p,Obs$x)$y)
```

This function is input to `modFit` which fits the model to the observations.

```
> print(system.time(
+ P      <- modFit(f=Residuals,p=c(0.1,1))
+ ))
```

```

user  system elapsed
 0      0          0

```

We can estimate and print the summary of fit

```
> sP    <- summary(P)
> sP
```

Parameters:

```
      Estimate Std. Error t value Pr(>|t|)
[1,]  0.14542    0.01564   9.296 0.000242 ***
[2,] 49.05292   17.91196   2.739 0.040862 *
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 0.01278 on 5 degrees of freedom

Parameter correlation:

```
      [,1] [,2]
[1,] 1.0000 0.8926
[2,] 0.8926 1.0000
```

We also plot the residual sum of squares, the residuals and the best-fit model

```
> x      <-0:375

> par(mfrow=c(2,2))
> plot(P,mfrow=NULL)
> plot(Obs,pch=16,cex=2,xlim=c(0,400),ylim=c(0,0.15),
+      xlab="mg COD/l",ylab="1/hr",main="best-fit")
> lines(Model(P$par,x))
> par(mfrow=c(1,1))
```

Finally, we run an MCMC analysis. The -scaled- parameter covariances returned from the `summary` function are used as estimate of the proposal covariances (`jump`). Scaling is as in (Gelman, Varlin, Stern, and Rubin 2004).

For the initial model variance (`var0`) we use the residual mean squares also returned by the `summary` function. We give equal weight to prior and modeled mean squares (`wvar0=1`)

The MCMC method adopted here is the metropolis-hastings algorithm; the MCMC is run for 3000 steps; we use the best-fit parameter set (`P$par`) to initiate the chain (`p`). A lower bound (0) is imposed on the parameters (`lower`).

```
> Covar    <- sP$cov.scaled * 2.4^2/2
> s2prior  <- sP$modVariance
> print(system.time(
+ MCMC <- modMCMC(f=Residuals,p=P$par,jump=Covar,niter=3000,
+               var0=s2prior,wvar0=1,lower=c(0,0))
+ ))
```

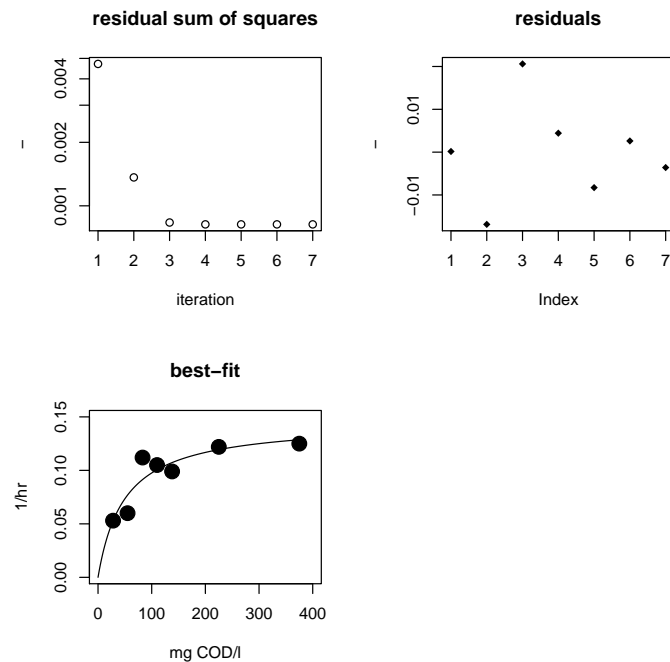


Figure 3: Fit diagnostics of the Monod function - see text for R-code

```
number of accepted runs: 1063 out of 3000 (35.43333%)
  user  system elapsed
 1.77    0.00    1.83
```

The plotted results demonstrate (near-) convergence of the chain.

```
> plot(MCMC, Full=TRUE)
```

The posterior distribution of the parameters, the sum of squares and the model's error standard deviation.

```
> hist(MCMC, Full=TRUE, col="darkblue")
```

The pairs plot shows the relationship between the two parameters

```
> pairs(MCMC)
```

The parameter correlation and covariances from the MCMC results can be calculated and compared with the results obtained by the fitting algorithm.

```
> cor(MCMC$pars)
```

```
      p1      p2
p1 1.000000 0.886937
p2 0.886937 1.000000
```

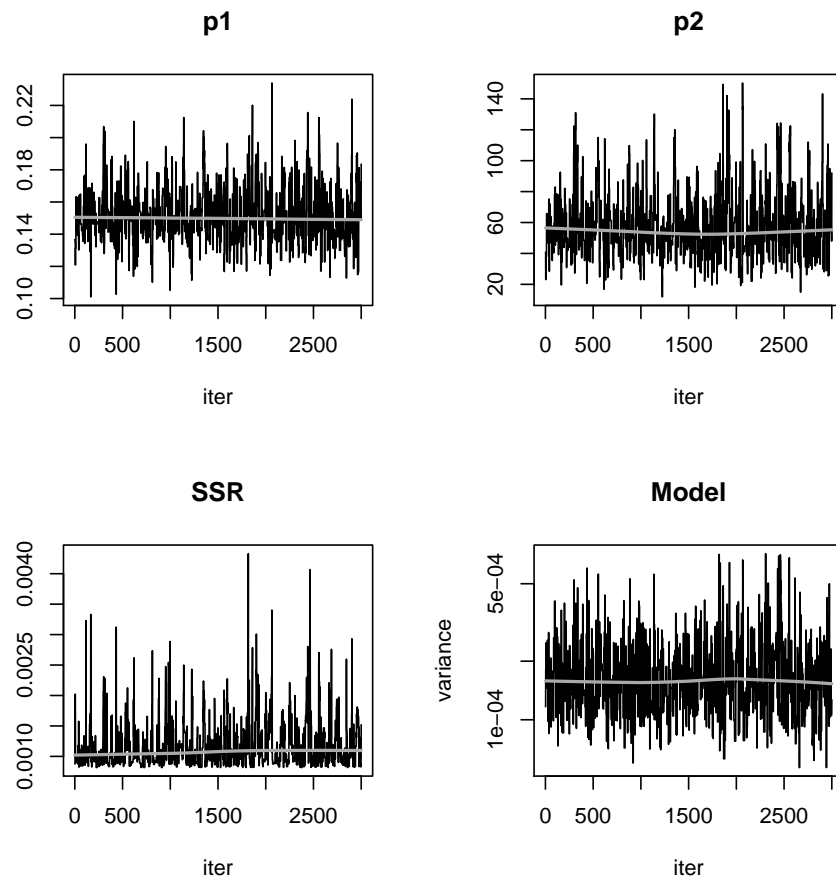



Figure 4: The mcmc - see text for R-code

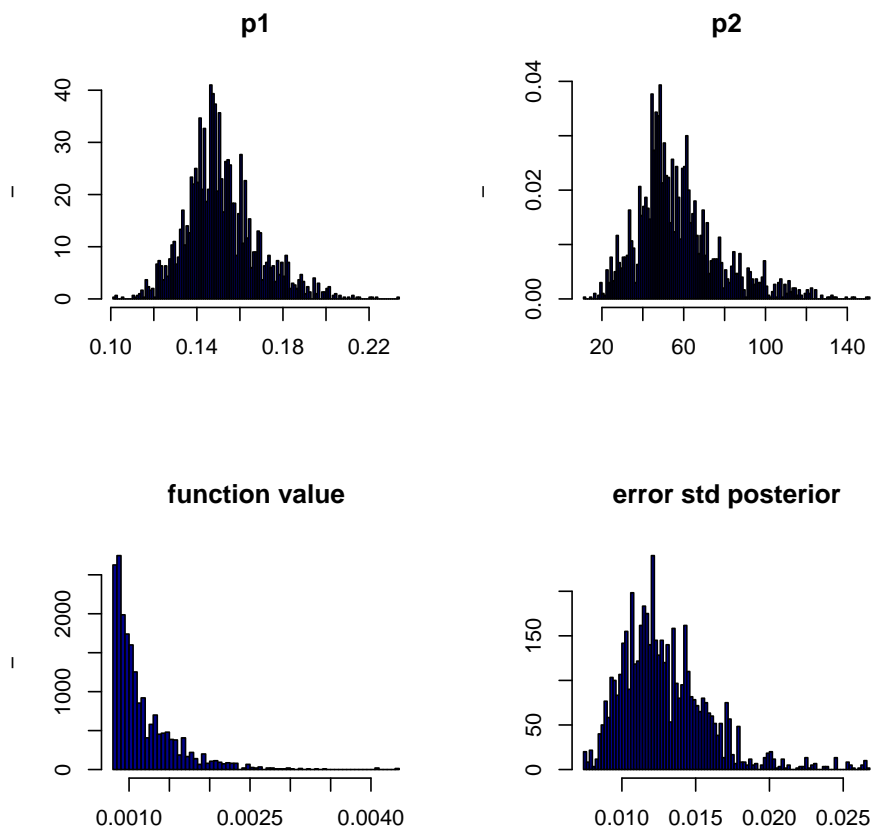


Figure 5: Hist plot - see text for R-code




Figure 6: Pairs plot - see text for R-code

```
> cov(MCMC$pars)
```

```
      p1      p2
p1 0.0002809622 0.3024043
p2 0.3024042765 413.7541673
```

```
> sP$cov.scaled
```

```
      [,1]      [,2]
[1,] 0.0002447075 0.2501157
[2,] 0.2501157147 320.8381590
```

The Raftery and Lewis's diagnostic from package **coda** gives more information on the number of runs that is actually needed. First the MCMC results need to be converted to an object of type `mcmc`, as used in **coda**.

```
> MC <- as.mcmc(MCMC$pars)
> raftery.diag(MC)
```

```
Quantile (q) = 0.025
Accuracy (r) = +/- 0.005
Probability (s) = 0.95
```

You need a sample size of at least 3746 with these values of q, r and s

Also interesting is function `cumuplot` from **coda**:

```
> cumuplot(MC)
```

The predictive posterior distribution of the model is easily estimated by running function `sensRange`, using a randomly selected subset of the parameters in the chain (`MCMC$pars`); we use the default of 100 parameter combinations.

```
> sR<-sensRange(parInput=MCMC$pars,func=Model,x=1:375)
```

The distribution is plotted and the data added to the plot:

```
> plot(summary(sR),quant=TRUE)
> points(Obs)
```

By toggling on covariance adaptation (`updatecov` and delayed rejection (`ntrydr`), the acceptance rate is increased:

```
> print(system.time(
+ MCMC2 <- modMCMC(f=Residuals,p=P$par,jump=Covar,niter=3000, ntrydr=3,
+                 var0=s2prior,wvar0=1,updatecov=100,lower=c(0,0))
+ ))
```




Figure 7: Cumulative quantile plot - see text for R-code




Figure 8: Predictive envelopes of the model - see text for R-code

```
number of accepted runs: 2425 out of 3000 (80.83333%)
```


```
  user  system elapsed
  4.35    0.00    4.34
```

```
> MCMC2$count
```

```
dr_steps      Alfasteps  num_accepted num_covupdate
      3100           13824           2425           30
```

4. A simple chemical model

This is an example from (Haario *et al.* 2006). We fit two parameters that describe the dynamics in the following reversible chemical reaction:



Here k_1 is the forward, k_2 the backward rate coefficient.

The ODE system is written as:

$$\begin{aligned}\frac{dA}{dt} &= -k_1 \cdot A + k_2 \cdot B \\ \frac{dB}{dt} &= +k_1 \cdot A - k_2 \cdot B\end{aligned}$$

with initial values $A_0 = 1$, $B_0 = 0$.

The analytical solution for this system of differential equations is given in (Haario *et al.* 2006).

First a function is defined that takes as input the parameters and that returns the values of the concentrations A and B, at selected output times.

```
> Reaction <- function (k, times)
+ {
+   fac <- k[1]/(k[1]+k[2])
+   A    <- fac + (1-fac)*exp(-(k[1]+k[2])*times)
+   return(data.frame(t=times,A=A))
+ }
```

All the concentrations were measured at the time the equilibrium was already reached. The data are the following:

```
> Data <- data.frame(
+   times = c(2, 4, 6, 8, 10),
+   A      = c(0.661, 0.668, 0.663, 0.682, 0.650))
> Data
```

	times	A
1	2	0.661
2	4	0.668
3	6	0.663
4	8	0.682
5	10	0.650

We need parameter priors to prevent the model parameters from drifting to infinite values. The prior is taken to be a broad Gaussian distribution with mean (2,4) and standard deviation = 200 for both.

The prior function returns the weighted sum of squared residuals of the parameter values with the expected value.

```
> Prior <- function(p)
+   return( sum(((p-c(2,4))/200)^2 ))
```

First the model is fitted to the data; we restrict the parameter values to be in the interval [0,1].

```
> residual <- function(k) return(Data$A - Reaction(k,Data$times)$A)
> Fit <- modFit(p=c(k1=0.5,k2=0.5),f=residual,lower=c(0,0),upper=c(1,1))
> (sF <- summary(Fit))
```

Parameters:

	Estimate	Std. Error	t value	Pr(> t)
k1	1.0000	0.3944	2.536	0.0850 .
k2	0.5123	0.1928	2.657	0.0765 .

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.01707 on 3 degrees of freedom

Parameter correlation:

	k1	k2
k1	1.000	0.996
k2	0.996	1.000

The residual error of the fit is used as initial model variance, the scaled covariance matrix of the fit is used as the proposal distribution (to generate new parameter values). As the covariance matrix is nearly singular this is not a very good approximation. The initial MCMC method, using the Metropolis-Hastings method does not converge. The MCMC is initiated with the best-fit parameters; the parameters are restricted to be positive numbers (**lower**).

```
> mse <- sF$modVariance
> Cov <- sF$cov.scaled * 2.4^2/2
> print(system.time(
+   MCMC <- modMCMC(f=residual, p=Fit$par, jump=Cov, lower=c(0,0),
+   +   var0=mse, wvar0=1, prior=Prior, niter=5000)
+ ))
```

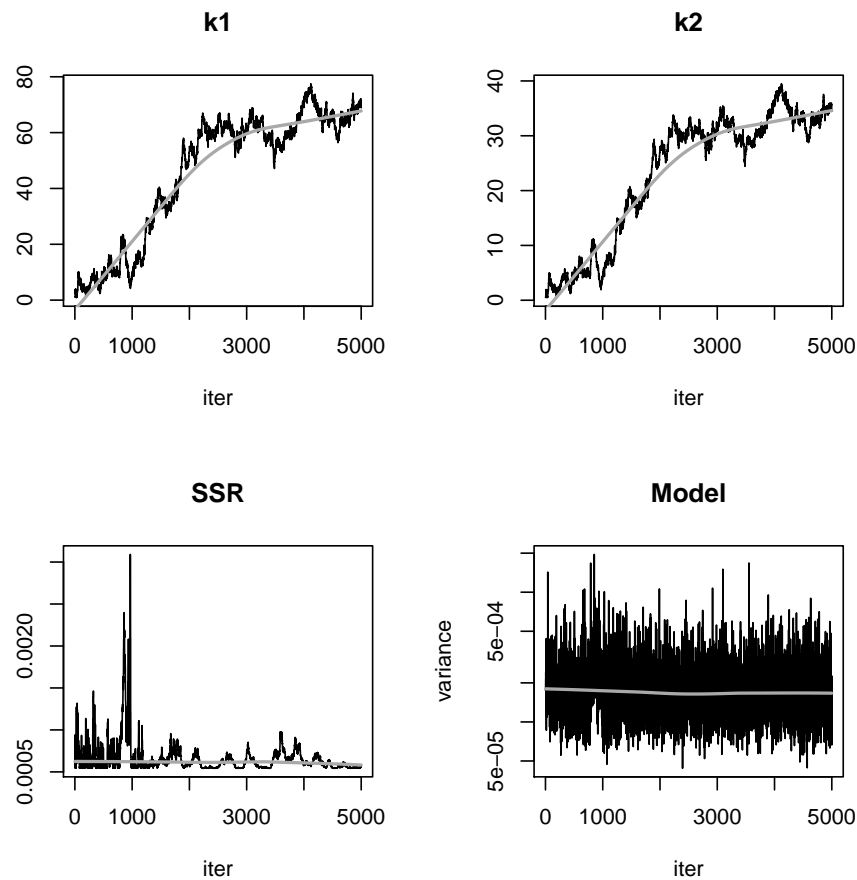



Figure 9: Metropolis-Hastings MCMC of the chemical model - see text for R-code

```
number of accepted runs: 4785 out of 5000 (95.7%)
  user  system elapsed
 2.96   0.00   2.97
```

The number of accepted runs is much too high, and indeed the MCMC has not at all converged...

```
> plot(MCMC, Full=TRUE)
```

Better convergence is achieved by the adaptive metropolis, updating the proposal every 100 runs

```
> MCMC2<- modMCMC(f=residual, p=Fit$par, jump=Cov, updatecov=100, lower=c(0,0),
+               var0=mse, wvar0=1, prior=Prior, niter=5000) #
```

```
number of accepted runs: 1486 out of 5000 (29.72%)
```

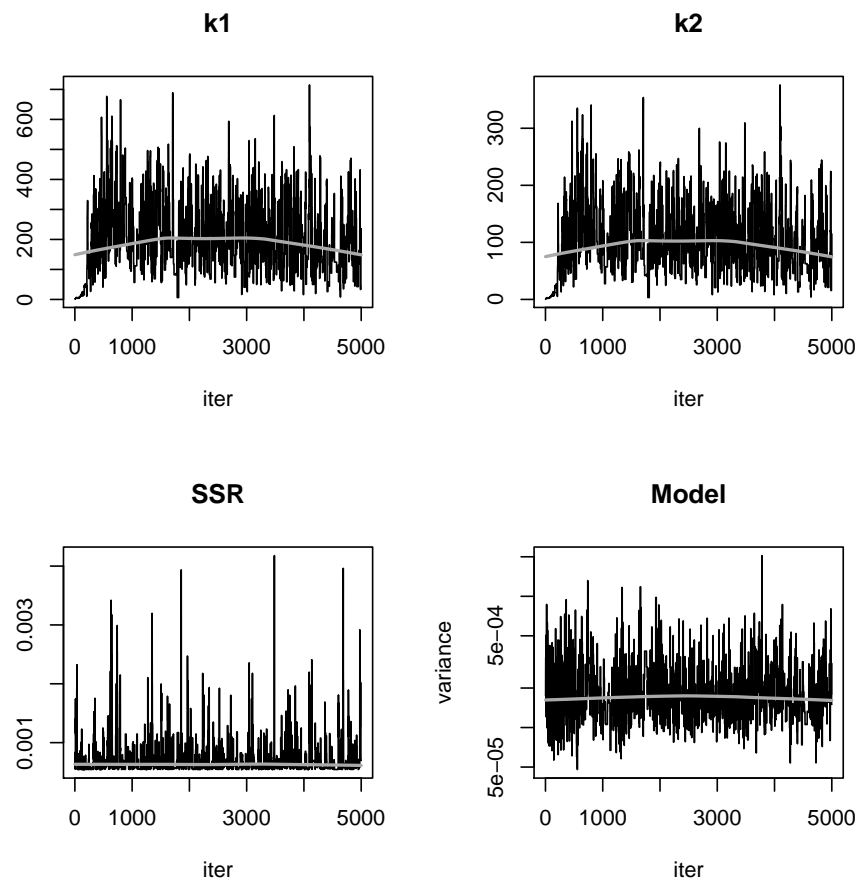


Figure 10: Adaptive Metropolis MCMC of the chemical model - see text for R-code




Figure 11: Pairs plot of the Adaptive Metropolis MCMC of the chemical model - see text for R-code

```
> plot(MCMC2,Full=TRUE)
```

The correlation between the two parameters is clear:

```
> pairs(MCMC2)
```

5. Oxygen in the sediment

5.1. the model

This is a simple model of oxygen, diffusing along a spatial gradient, with imposed upper boundary concentration oxygen is consumed at maximal fixed rate, and including a monod limitation.

The constitutive equations are:

$$\frac{\partial O_2}{\partial t} = -\frac{\partial Flux}{\partial x} - cons \cdot \frac{O_2}{O_2 + k_s}$$

$$Flux = -D \cdot \frac{\partial O_2}{\partial x}$$

$$O_2(x = 0) = upO2$$

```
> par(mfrow=c(2,2))
> require(FME)
```

First the model parameters are defined...

```
> pars <- c(upO2=360, # concentration at upper boundary, mmolO2/m3
+          cons=80,   # consumption rate, mmolO2/m3/day
+          ks=1,      # O2 half-saturation ct, mmolO2/m3
+          D=1)       # diffusion coefficient, cm2/d
```

Next the sediment is vertically subdivided into 100 grid cells, each 0.05 cm thick.

```
> n <- 100 # nr grid points
> dx <- 0.05 #cm
> dX <- c(dx/2,rep(dx,n-1),dx/2) # dispersion distances; half dx near boundaries
> X <- seq(dx/2,len=n,by=dx) # distance from upper interface at middle of box
```

The model function takes as input the parameter values and returns the steady-state condition of oxygen. Function `steady.band` from package **rootSolve** ((Soetaert 2008)) does this in a very efficient way (see (Soetaert and Herman 2009)).

```
> O2fun <- function(pars)
+ {
+   derivs<-function(t,O2,pars)
+   {
+     with (as.list(pars),{
+       Flux <- -D* diff(c(upO2,O2,O2[n]))/dX
+       dO2 <- -diff(Flux)/dx-cons*O2/(O2+ks)
+       return(list(dO2,UpFlux = Flux[1],LowFlux = Flux[n+1]))
+     })
+   }
+   # Solve the steady-state conditions of the model
+   ox <- steady.band(y=runif(n),func=derivs,parms=pars,nspec=1,positive=TRUE)
+   data.frame(X=X,O2=ox$y)
+ }
```

The model is run




Figure 12: The modeled oxygen profile - see text for R-code

```
> ox<-O2fun(pars)
```

and the results plotted...

```
> plot(ox$O2,ox$X,ylim=rev(range(X)),xlab="mmol/m3",
+      main="Oxygen", ylab="depth, cm",type="l",lwd=2)
```

5.2. Global sensitivity analysis : Sensitivity ranges

The sensitivity of the oxygen profile to parameter `cons`, the consumption rate is estimated. We assume a normally distributed parameter, with mean = 80 (`parMean`), and a variance=100 (`parCovar`). The model is run 100 times (`num`).

```
> print(system.time(
+ Sens2 <- sensRange(parms=pars,func=O2fun,dist="norm",
+                   num=100,parMean=c(cons=80),parCovar=100)
+ ))
```

```
user  system elapsed
0.99   0.00   0.99
```

The results can be plotted in two ways:

```
> par(mfrow=c(1,2))
> plot(Sens2,xyswap=TRUE,xlab= "O2",
+      ylab="depth, cm",main="Sensitivity runs")
> plot(summary(Sens2),xyswap=TRUE,xlab= "O2",
+      ylab="depth, cm",main="Sensitivity ranges")
> par(mfrow=c(1,1))
```




Figure 13: Results of the sensitivity run - left: all model runs, right: summary - see text for R-code

5.3. Local sensitivity analysis : Sensitivity functions

Local sensitivity analysis starts by calculating the sensitivity functions

```
> O2sens <- sensFun(func=O2fun,parms=pars)
```

The summary of these functions gives information about which parameters have the largest effect (univariate sensitivity):

```
> summary(O2sens)
```

	value	scale	L1	L2	Mean	Min	Max	N
upO2	360	360	7.0	0.88	7.0	1.0e+00	13.4176	100
cons	80	80	8.2	1.16	-8.2	-2.2e+01	-0.0084	100
ks	1	1	2.2	0.37	2.2	1.2e-04	9.6137	100
D	1	1	7.4	1.02	7.4	8.4e-03	19.9515	100

In bivariate sensitivity the pair-wise relationship and the correlation is estimated and/or plotted:

```
> pairs(O2sens)
```

```
> cor(O2sens[,-(1:2)])
```

	upO2	cons	ks	D
upO2	1.0000000	-0.9785990	0.8375806	0.9795154
cons	-0.9785990	1.0000000	-0.9320469	-0.9992101
ks	0.8375806	-0.9320469	1.0000000	0.9288055
D	0.9795154	-0.9992101	0.9288055	1.0000000




Figure 14: pairs plot - see text for R-code

Multivariate sensitivity is done by estimating the collinearity between parameter sets.

```
> Coll <- collin(O2sens)
> Coll
```

	upO2	cons	ks	D	N	collinearity
1	1	1	0	0	2	7.6
2	1	0	1	0	2	2.9
3	1	0	0	1	2	8.1
4	0	1	1	0	2	4.4
5	0	1	0	1	2	46.1
6	0	0	1	1	2	4.2
7	1	1	1	0	3	24.9
8	1	1	0	1	3	49.9
9	1	0	1	1	3	26.3
10	0	1	1	1	3	49.3
11	1	1	1	1	4	49.9

```
> plot(Coll, log="y")
```

5.4. Fitting the model to the data

Assume both the oxygen flux at the upper interface and a vertical profile of oxygen has been measured.

These are the data:

```
> O2dat <- data.frame(x=seq(0.1,3.5,by=0.1),
+   y = c(279,260,256,220,200,203,189,179,165,140,138,127,116,
+   109,92,87,78,72,62,55,49,43,35,32,27,20,15,15,10,8,5,3,2,1,0))
```




Figure 15: collinearity - see text for R-code

```
> O2depth <- cbind(name="O2",O2dat)          # oxygen versus depth
> O2flux  <- c(UpFlux=170)                   # measured flux
```

First a function is defined that returns only the required model output.

```
> O2fun2 <- function(pars)
+ {
+   derivs<-function(t,O2,pars)
+   {
+     with (as.list(pars),{
+       Flux <- -diff(c(upO2,O2,O2[n]))/dX
+       dO2  <- -diff(Flux)/dx-cons*O2/(O2+ks)
+       return(list(dO2,UpFlux = Flux[1],LowFlux = Flux[n+1]))
+     })
+   }
+
+   ox <- steady.band(y=runif(n),func=derivs,parms=pars,nspec=1,positive=TRUE)
+
+   list(data.frame(x=X,O2=ox$y),
+         UpFlux=ox$UpFlux)
+ }
```

The function used in the fitting algorithm returns an instance of type `modCost`. This is created by calling function `modCost` twice. First with the modeled oxygen profile, then with the modeled flux.

```
> Objective <- function (P)
+ {
```



```

+ pars[names(P)]<-P
+ mod02 <- 02fun2(pars)
+
+ # Model cost: first the oxygen profile
+ Cost <- modCost(obs=02depth,model=mod02[[1]],x="x",y="y")
+
+ # then the flux
+ modF1 <- c(UpFlux=mod02$UpFlux)
+ Cost <- modCost(obs=02flux,model=modF1,x=NULL,cost=Cost)
+
+ return(Cost)
+ }

```

We first estimate the identifiability of the parameters, given the data:

```

> print(system.time(
+ sF<-sensFun(Objective,parms=c(up02=360,cons=80,ks=1,D=1))
+ ))

```

```

      user  system elapsed
0.10      0.00      0.09

```

```

> collin(sF)

```

	up02	cons	ks	D	N	collinearity
1	1	1	0	0	2	8.0
2	1	0	1	0	2	3.0
3	1	0	0	1	2	1.0
4	0	1	1	0	2	4.2
5	0	1	0	1	2	1.1
6	0	0	1	1	2	1.2
7	1	1	1	0	3	14.3
8	1	1	0	1	3	10.7
9	1	0	1	1	3	4.7
10	0	1	1	1	3	6.3
11	1	1	1	1	4	14.3

The value of the full set is relatively large, and the oxygen diffusion coefficient is well-known, so it is left out of the fitting. The combination of the three remaining parameters has a low enough collinearity to enable automatic fitting. The parameters are constrained to be >0

```

> collin(sF,parset=c("up02","cons","ks"))

```

	up02	cons	ks	D	N	collinearity
1	1	1	1	0	3	14

```
> print(system.time(
+ Fit<-modFit(p=c(up02=360,cons=80,ks=1),
+               f=Objective,lower=c(0,0,0))
+             ))

      user  system elapsed
      0.93    0.00    0.92

> (SFit<-summary(Fit))

Parameters:
      Estimate Std. Error t value Pr(>|t|)
up02   292.936      2.104 139.259  <2e-16 ***
cons    49.683      2.360  21.054  <2e-16 ***
ks       1.295      1.353   0.957   0.345
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.401 on 33 degrees of freedom

Parameter correlation:
      up02   cons    ks
up02 1.0000 0.5795 0.2972
cons 0.5795 1.0000 0.9006
ks   0.2972 0.9006 1.0000
```

We next plot the residuals

```
> plot(Objective(Fit$par),xlab="depth",ylab="",main="residual",legpos="top")
```

and show the best-fit model

```
> pars[names(Fit$par)]<- Fit$par
> mod02 <- 02fun(pars)

> plot(02depth$y,02depth$x,ylim=rev(range(02depth$x)),pch=18,
+       main="Oxygen-fitted", xlab="mmol/m3",ylab="depth, cm")
> lines(mod02$02,mod02$X)
```

5.5. Run MCMC

We use the parameter covariances of previous fit to update parameters, while the mean squared residual of the fit is use as prior fo the model variance.

```
> Covar   <- SFit$cov.scaled * 2.4^2/4
> s2prior <- SFit$modVariance
```




Figure 16: residuals - see text for R-code




Figure 17: Best fit model - see text for R-code

We run an adaptive metropolis, making sure that ks does not become negative...

```
> print(system.time(
+ MCMC <- modMCMC(f=Objective,p=Fit$par,jump=Covar,niter=500,ntrydr=2,
+                 var0=s2prior,wvar0=1,updatecov=100,lower=c(NA,NA,0))
+ ))
```

number of accepted runs: 401 out of 500 (80.2%)

```
  user  system elapsed
12.53   0.02   12.61
```

```
> MCMC$count
```

```
      dr_steps      Alfasteps  num_accepted num_covupdate
          229           687           401           5
```

Plotting the results is similar to previous cases.

```
> plot(MCMC,Full=TRUE)
```

```
> hist(MCMC,Full=TRUE)
```

```
> pairs(MCMC,Full=TRUE)
```

or summaries can be created:

```
> summary(MCMC)
```

```
      up02      cons      ks      sig
mean 295.120874  69.84969 16.66351517 1602.845769
sd    8.621489  36.37423 27.43107976 7107.847836
min  269.343451  37.19269  0.01983037  1.448531
max  342.116114 224.05698 131.33231446 94072.730764
q025 291.726850  49.98458  1.63636846  13.020102
q050 294.021329  53.00745  4.13958583  54.054699
q075 297.167247  71.25788 17.56681409 341.060071
```

```
> cor(MCMC$pars)
```

```
      up02      cons      ks
up02 1.0000000 0.5986622 0.5812072
cons 0.5986622 1.0000000 0.9951063
ks    0.5812072 0.9951063 1.0000000
```

Note: we pass to `sensRange` the full parameter vector (`parms`) and the parameters sampled during the MCMC (`parInput`).

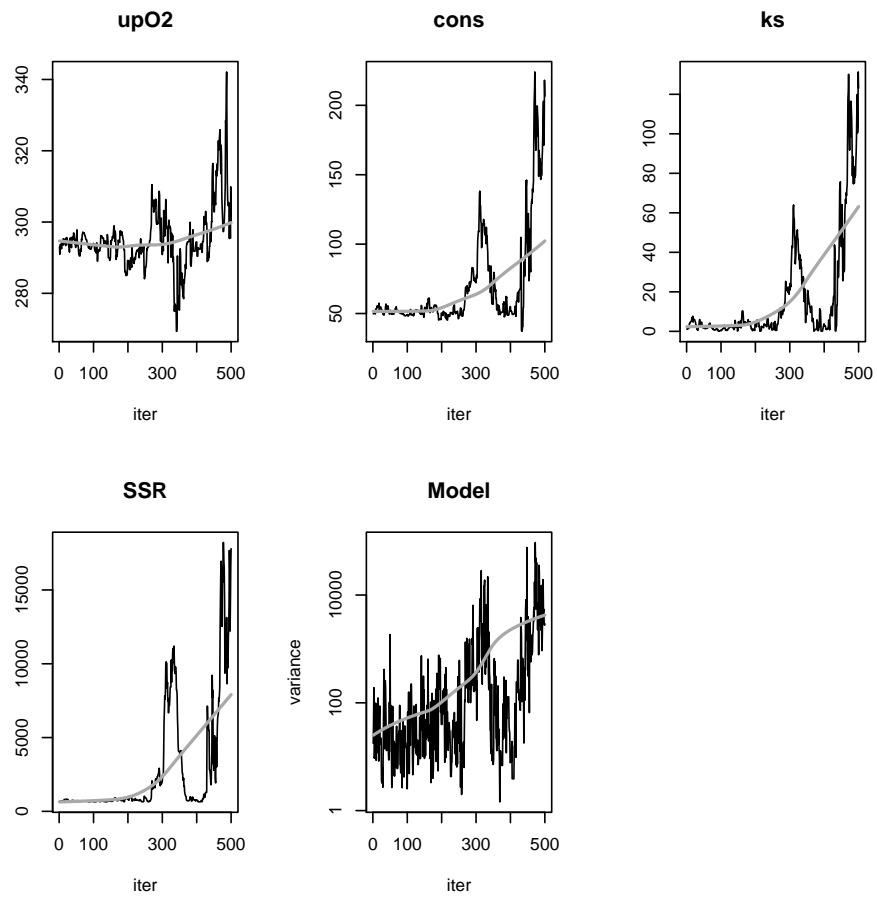


Figure 18: MCMC plot results - see text for R-code

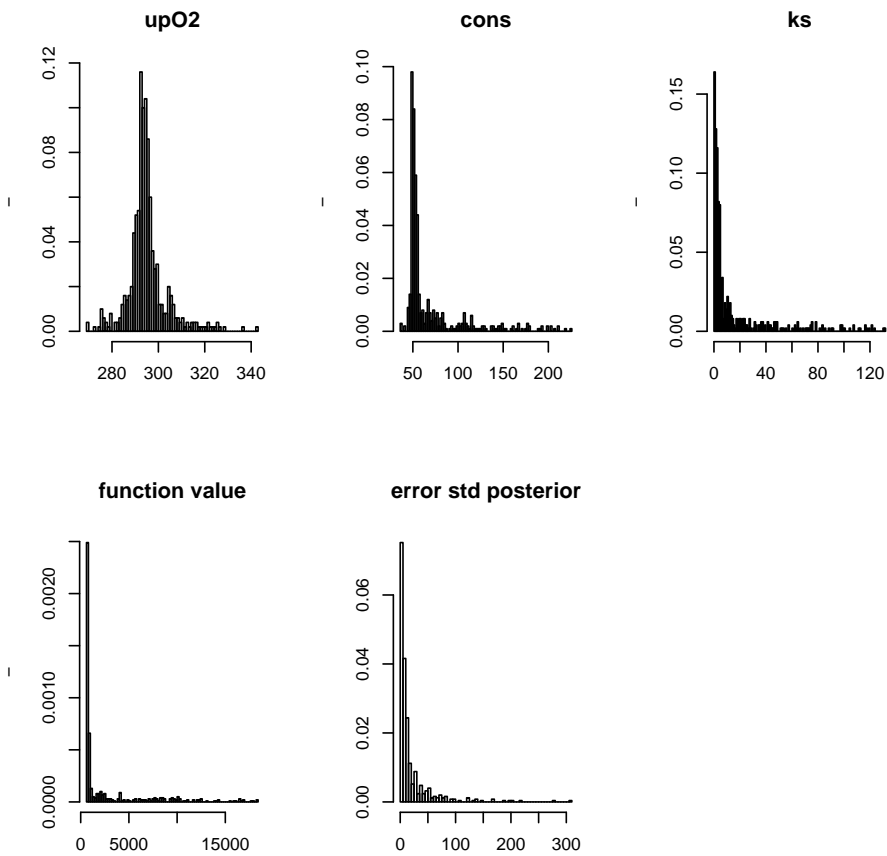


Figure 19: MCMC histogram results - see text for R-code




Figure 20: MCMC pairs plot - see text for R-code




Figure 21: MCMC range plot - see text for R-code


```
> plot(summary(sensRange(parms=pars,parInput=MCMC$par,f=02fun,num=500)),xyswap=TRUE)
> points(02depth$y,02depth$x)
```

6. finally

This vignette is a Sweave (Leisch 2002) translation of part of the **FME** examples.

References

- Gelman A, Varlin JB, Stern HS, Rubin DB (2004). *Bayesian Data Analysis, second edition*. Chapman and Hall / CRC, Boca Raton.
- Haario H, Laine M, Mira A, Saksman E (2006). “DRAM: efficient adaptive MCMC.” *Statistical Computing*, **16**, 339–354.
- Laine M (2008). *Adaptive MCMC methods with applications in environmental and geophysical models*. Finnish meteorological institute contributions n0 69 -ISBN 978-951-697-662-7.
- Leisch F (2002). “Sweave: Dynamic Generation of Statistical Reports Using Literate Data Analysis.” In W Härdle, B Rönz (eds.), “Compstat 2002 - Proceedings in Computational Statistics,” pp. 575–580. Physica Verlag, Heidelberg. ISBN 3-7908-1517-9, URL <http://www.stat.uni-muenchen.de/~leisch/Sweave>.
- Soetaert K (2008). *rootSolve: Nonlinear root finding, equilibrium and steady-state analysis of ordinary differential equations*. R package version 1.2.
- Soetaert K, Herman PMJ (2009). *A Practical Guide to Ecological Modelling. Using Ras a Simulation Platform*. Springer. ISBN 978-1-4020-8623-6.

Affiliation:

Karline Soetaert
 Centre for Estuarine and Marine Ecology (CEME)
 Netherlands Institute of Ecology (NIOO)
 4401 NT Yerseke, Netherlands
 E-mail: k.soetaert@nioo.knaw.nl
 URL: <http://www.nioo.knaw.nl/ppages/ksoetaert>